
MGT-415: Data Science in Practice

Project Notebook
May 10, 2019

Do Carmo Maria Genc Murat Nyambuu Lkham Trichilo Giulio Xi Fan



ÉCOLE POLYTECHNIQUE
FÉDÉRALE DE LAUSANNE

I PYTHON DEPENDENCIES

```
In [1]: ! pip install missingno #missing data
        ! pip install inblearn #over/undersampling

Requirement already satisfied: missingno in
/home/zarattras/anaconda3/lib/python3.7/site-packages (0.4.1)
Requirement already satisfied: scipy in /home/zarattras/anaconda3/lib/python3.7/site-
packages (from missingno) (1.2.1)
Requirement already satisfied: matplotlib in
/home/zarattras/anaconda3/lib/python3.7/site-packages (from missingno) (3.0.3)
Requirement already satisfied: seaborn in /home/zarattras/anaconda3/lib/python3.7/site-
packages (from missingno) (0.9.0)
Requirement already satisfied: numpy in /home/zarattras/anaconda3/lib/python3.7/site-
packages (from missingno) (1.16.3)
Requirement already satisfied: cycler>=0.10 in
/home/zarattras/anaconda3/lib/python3.7/site-packages (from matplotlib->missingno)
(0.10.0)
Requirement already satisfied: kiwisolver>=1.0.1 in
/home/zarattras/anaconda3/lib/python3.7/site-packages (from matplotlib->missingno)
(1.0.1)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in
/home/zarattras/anaconda3/lib/python3.7/site-packages (from matplotlib->missingno)
(2.3.1)
Requirement already satisfied: python-dateutil>=2.1 in
/home/zarattras/anaconda3/lib/python3.7/site-packages (from matplotlib->missingno)
(2.8.0)
Requirement already satisfied: pandas>=0.15.2 in
/home/zarattras/anaconda3/lib/python3.7/site-packages (from seaborn->missingno)
(0.24.2)
Requirement already satisfied: six in /home/zarattras/anaconda3/lib/python3.7/site-
packages (from cycler>=0.10->matplotlib->missingno) (1.12.0)
Requirement already satisfied: setuptools in
/home/zarattras/anaconda3/lib/python3.7/site-packages (from
kiwisolver>=1.0.1->matplotlib->missingno) (40.8.0)
Requirement already satisfied: pytz>=2011k in
/home/zarattras/anaconda3/lib/python3.7/site-packages (from
pandas>=0.15.2->seaborn->missingno) (2018.9)
Collecting inblearn
  Could not find a version that satisfies the requirement inblearn (fromversions: )
No matching distribution found for inblearn
```

II DATASET ANALYSIS

```
In [2]: import numpy as np
        import pandas as pd
        import matplotlib.pyplot as plt
        import seaborn as sns
        import missingno as msno

        #####

        import numpy as np
```

```

import scipy.stats as sc
import pandas as pd

import seaborn as sns
import matplotlib as mpl

import matplotlib.pyplot as plt
from IPython.display import display

import statsmodels.api as sm

from sklearn.model_selection import train_test_split
from statsmodels.stats import outliers_influence as oi

from sklearn.linear_model import LogisticRegression
from sklearn.linear_model import LinearRegression

from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.metrics import log_loss
from sklearn.metrics import mean_squared_error
from sklearn.feature_selection import RFE
from sklearn import metrics

from sklearn.ensemble import IsolationForest
from sklearn.model_selection import cross_val_predict
#####

from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn import metrics

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
from sklearn import metrics

from matplotlib.colors import ListedColormap
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.datasets import make_moons, make_circles, make_classification
#from sklearn.neural_network import MLPClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.gaussian_process import GaussianProcessClassifier
from sklearn.gaussian_process.kernels import RBF
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis
from sklearn.calibration import CalibratedClassifierCV

#####

```

```

#from sklearn.datasets import load_iris
from sklearn import preprocessing

from scipy import interp
from sklearn import datasets, neighbors
from sklearn.metrics import auc, roc_curve
from sklearn.model_selection import StratifiedKFold

from imblearn.over_sampling import ADASYN, SMOTE, RandomOverSampler, SMOTENC
from imblearn.combine import SMOTEENN
from imblearn.under_sampling import ClusterCentroids,
RandomUnderSampler, EditedNearestNeighbours

from imblearn.pipeline import make_pipeline
from imblearn import FunctionSampler
#####
import warnings
warnings.filterwarnings("ignore")
pd.set_option('display.max_colwidth', 250)
sns.set()

```

II Basic Info

```

In [3]: dsdata = pd.read_excel('Data/full_dataset.xlsx')
dsdata = dsdata.replace('unknown', np.nan)
dsdata = dsdata.replace('nonexistent', np.nan)

dsdata.head()

```

```

Out[3]:   age      job  marital  education default housing loan   contact month \
0   56  housemaid  married   basic.4y      no      no  no  telephone  may
1   57  services  married  high.school   NaN      no  no  telephone  may
2   37  services  married  high.school      no     yes  no  telephone  may
3   40   admin.  married   basic.6y      no      no  no  telephone  may
4   56  services  married  high.school      no      no  yes  telephone  may

```

```

   day_of_week  ...  campaign  pdays  previous  poutcome emp.var.rate \
0         mon  ...         1     999         0        NaN         1.1
1         mon  ...         1     999         0        NaN         1.1
2         mon  ...         1     999         0        NaN         1.1
3         mon  ...         1     999         0        NaN         1.1
4         mon  ...         1     999         0        NaN         1.1

```

```

   cons.price.idx  cons.conf.idx  euribor3m  nr.employed  y
0         93.994        -36.4      4.857      5191.0  no
1         93.994        -36.4      4.857      5191.0  no
2         93.994        -36.4      4.857      5191.0  no
3         93.994        -36.4      4.857      5191.0  no
4         93.994        -36.4      4.857      5191.0  no

```

[5 rows x 21 columns]

```

In [4]: dsdata.dtypes

```

```

Out[4]: age                int64
job                object

```

```

marital      object
education    object
default      object
housing      object
loan         object
contact      object
month        object
day_of_week  object
duration     int64
campaign     int64
pdays       int64
previous     int64
poutcome     object
emp.var.rate float64
cons.price.idx float64
cons.conf.idx float64
euribor3m    float64
nr.employed  float64
y            object
dtype: object

```

II Input variables:

a. *bank client data:* 1 - age (numeric)

2 - job : type of job (categorical: 'admin.', 'blue-collar', 'entrepreneur', 'housemaid', 'management', 'retired', 'self-employed', 'services', 'student', 'technician', 'unemployed', 'unknown')

3 - marital : marital status (categorical: 'divorced', 'married', 'single', 'unknown'; note: 'divorced' means divorced or widowed)

4 - education (categorical: 'basic.4y', 'basic.6y', 'basic.9y', 'high.school', 'illiterate', 'professional.course', 'university.degree', 'unknown')

5 - default: has credit in default? (categorical: 'no', 'yes', 'unknown')

6 - housing: has housing loan? (categorical: 'no', 'yes', 'unknown')

7 - loan: has personal loan? (categorical: 'no', 'yes', 'unknown')

1. related with the last contact of the current campaign:

8 - contact: contact communication type (categorical: 'cellular', 'telephone')

9 - month: last contact month of year (categorical: 'jan', 'feb', 'mar', ..., 'nov', 'dec')

10 - day_of_week: last contact day of the week (categorical: 'mon', 'tue', 'wed', 'thu', 'fri')

11 - duration: last contact duration, in seconds (numeric). Important note: this attribute highly affects the output target (e.g., if duration=0 then y='no'). Yet, the duration is not known before a call is performed. Also, after the end of the call y is obviously known. Thus, this input should only be included for benchmark purposes and should be discarded if the intention is to have a realistic predictive model.

2. other attributes:

12 - campaign: number of contacts performed during this campaign and for this client (numeric, includes last contact)

13 - pdays: number of days that passed by after the client was last contacted from a previous campaign (numeric; 999 means client was not previously contacted)

14 - previous: number of contacts performed before this campaign and for this client (numeric)

15 - poutcome: outcome of the previous marketing campaign (categorical: 'failure', 'nonexistent', 'success')

3. social and economic context attributes

- 16 - emp.var.rate: employment variation rate - quarterly indicator (numeric)
- 17 - cons.price.idx: consumer price index - monthly indicator (numeric)
- 18 - cons.conf.idx: consumer confidence index - monthly indicator (numeric)
- 19 - euribor3m: euribor 3 month rate - daily indicator (numeric)
- 20 - nr.employed: number of employees - quarterly indicator (numeric)

4. Output variable (desired target):

- 21 - y - has the client subscribed a term deposit? (binary: 'yes','no')

II Variable Description

In [5]: *#describe the "Object" type elements*

```
dsdata.describe(include=['O'])
```

```
Out[5]:
```

	job	marital	education	default	housing	loan	contact	\
count	40858	41108	39457	32591	40198	40198	41188	
unique	11	3	7	2	2	2	2	
top	admin.	married	university.degree	no	yes	no	cellular	
freq	10422	24928	12168	32588	21576	33950	26144	

	month	day_of_week	poutcome	y
count	41188	41188	5625	41188
unique	10	5	2	2
top	may	thu	failure	no
freq	13769	8623	4252	36548

In [6]: *#describe the "numerical" type elements*

```
dsdata.describe()
```

```
Out[6]:
```

	age	duration	campaign	pdays	previous	\
count	41188.00000	41188.00000	41188.00000	41188.00000	41188.00000	
mean	40.02406	258.285010	2.567593	962.475454	0.172963	
std	10.42125	259.279249	2.770014	186.910907	0.494901	
min	17.00000	0.000000	1.000000	0.000000	0.000000	
25%	32.00000	102.000000	1.000000	999.000000	0.000000	
50%	38.00000	180.000000	2.000000	999.000000	0.000000	
75%	47.00000	319.000000	3.000000	999.000000	0.000000	
max	98.00000	4918.000000	56.000000	999.000000	7.000000	

	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	nr.employed
count	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000
mean	0.081886	93.575664	-40.502600	3.621291	5167.035911
std	1.570960	0.578840	4.628198	1.734447	72.251528
min	-3.400000	92.201000	-50.800000	0.634000	4963.600000
25%	-1.800000	93.075000	-42.700000	1.344000	5099.100000
50%	1.100000	93.749000	-41.800000	4.857000	5191.000000
75%	1.400000	93.994000	-36.400000	4.961000	5228.100000
max	1.400000	94.767000	-26.900000	5.045000	5228.100000




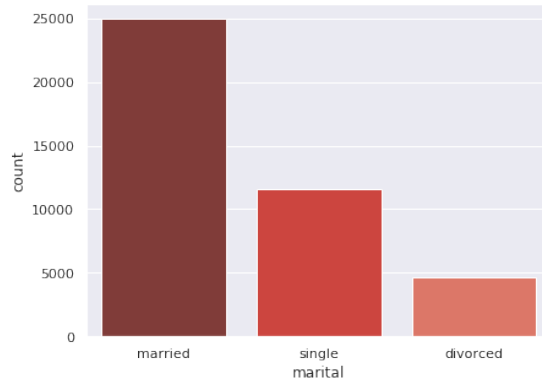


In [7]: `dsdata['y'].value_counts()`

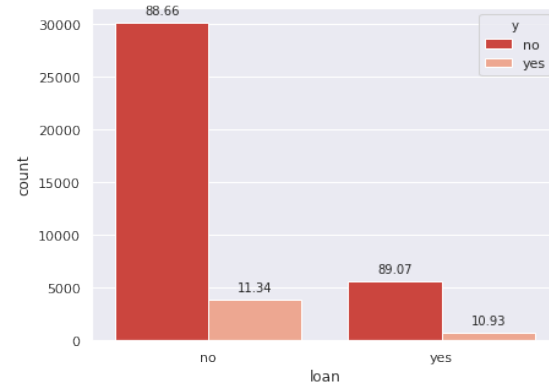
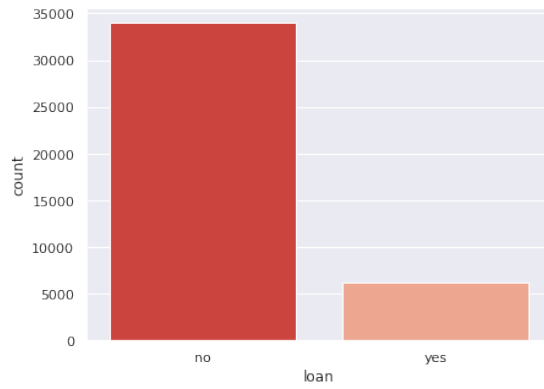

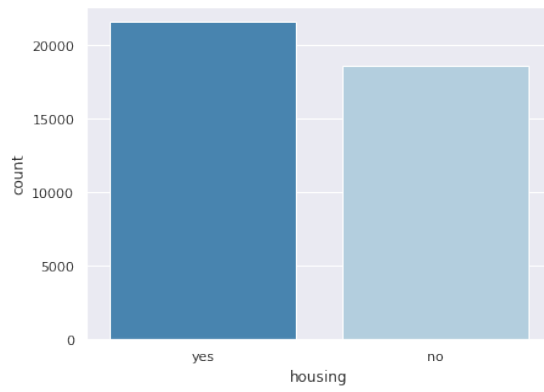
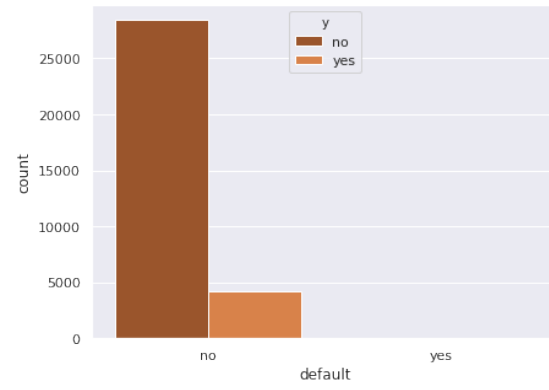

```
Out[7]: no      36548
        yes      4640
        Name: y, dtype: int64
```

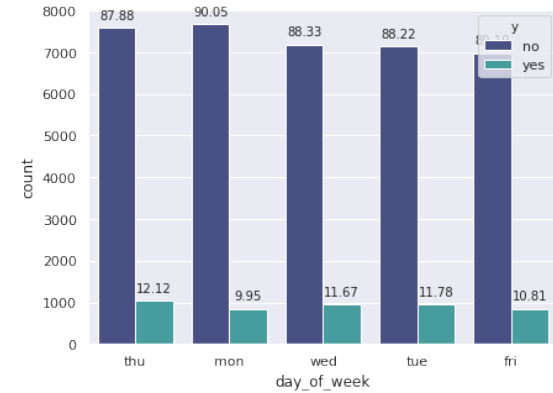



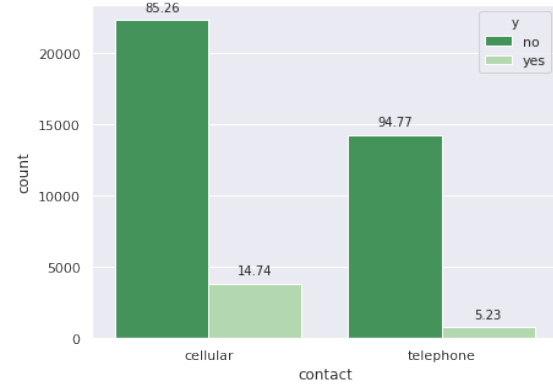
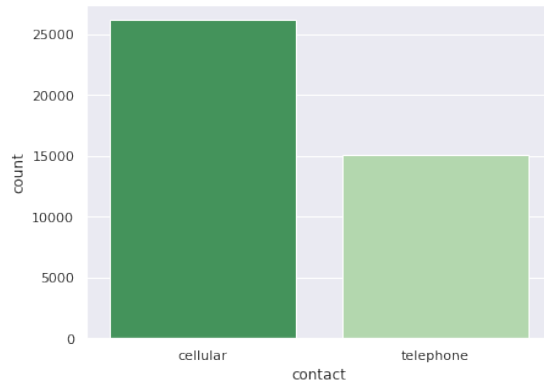
```
In [ ]:
```


II Variable Distribution

```
In [8]: # Barplots for categorical (object) variables and comparison between yes/no in
deposit(y)
colors = ['Blues_d', 'Reds_d', 'Greens_d', 'Oranges_d', 'Blues_r', 'Reds_r', 'Greens_r',
'Oranges_r', 'mako', 'Reds_d', 'BuPu_r']
color_coef=0
total = float(len(dsdata))
for column in dsdata.columns:
    if dsdata[column].dtype == "O":
        if column=='job' or column=='education':
            x = dsdata[column].value_counts()
            fig,ax=plt.subplots(1,2,figsize=(15,5))
            t = sns.countplot(x=column, data=dsdata,
palette=colors[color_coef],order=dsdata[column].value_counts().index, ax=ax[0])
            t.set_xticklabels(t.get_xticklabels(), rotation=90)
            g = sns.countplot(x=dsdata[column], hue=dsdata['y'],
palette=colors[color_coef], order=dsdata[column].value_counts().index, ax=ax[1])
            i=0
            for p in g.patches:
                length = len(g.patches)/2
                total = x[i]
                g.annotate(format(p.get_height()/total*100, '.2f'), (p.get_x() +
p.get_width() / 2., p.get_height()), ha = 'center', va = 'center', xytext = (0, 10),
textcoords = 'offset points')
                i +=1
            if i==length:
                i=0
            g.set_xticklabels(g.get_xticklabels(), rotation=90)
        else:
            x = dsdata[column].value_counts()
            fig,ax=plt.subplots(1,2,figsize=(15,5))
            t = sns.countplot(x=column, data=dsdata, palette=colors[color_coef],
order=dsdata[column].value_counts().index, ax=ax[0])
            g = sns.countplot(x=dsdata[column], hue=dsdata['y'],
palette=colors[color_coef],order=dsdata[column].value_counts().index, ax=ax[1])
            if column=='default' or column=='y':
                pass
            else:
                i=0
                for p in g.patches:
                    length = len(g.patches)/2
                    total = x[i]
                    g.annotate(format(p.get_height()/total*100, '.2f'), (p.get_x() +
p.get_width() / 2., p.get_height()), ha = 'center', va = 'center', xytext = (0, 10),
textcoords = 'offset points')
                    i +=1
                if i==length:
                    i=0
            color_coef +=1
```









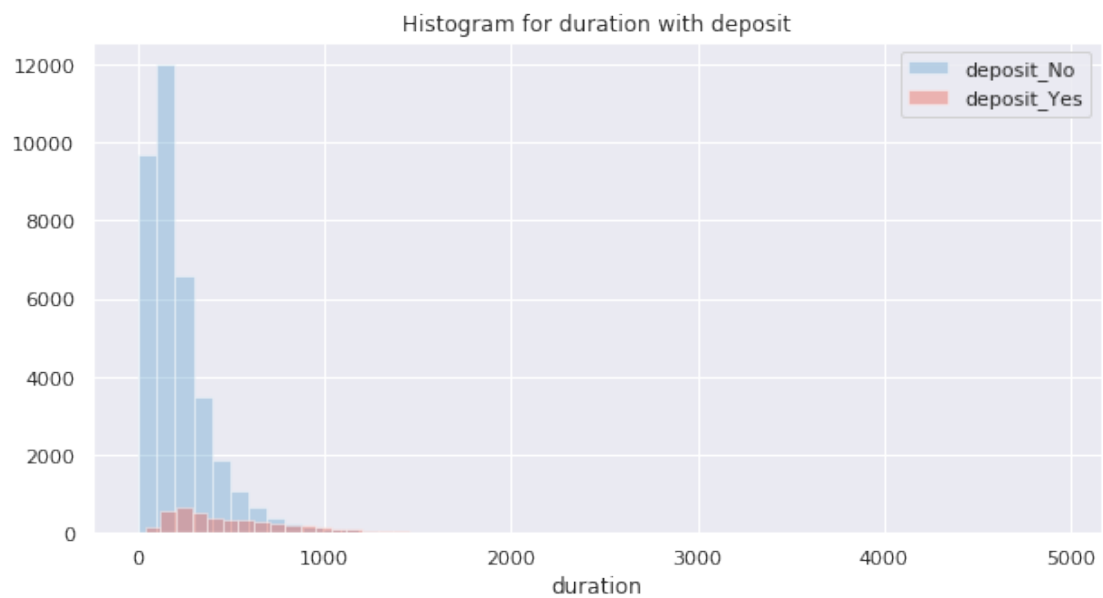




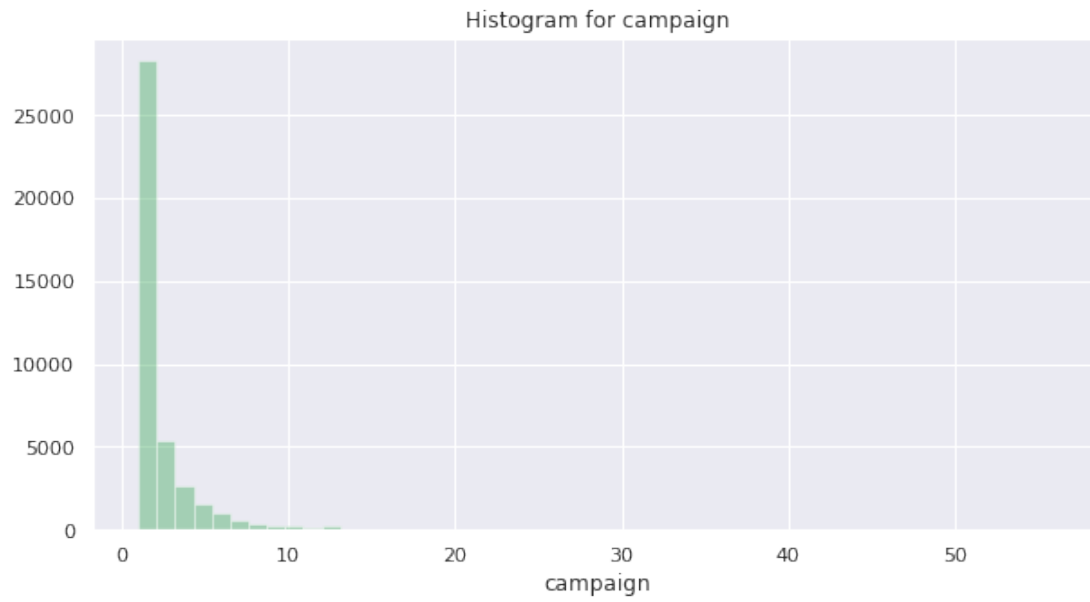
```
In [9]: def histogram(variable):
    plt.figure(figsize=(10, 5))
    plt.title("Histogram for {}".format(variable))
    ax = sns.distplot(dsdata[variable], color=sns.color_palette("RdYlGn_r")[0], kde=False)

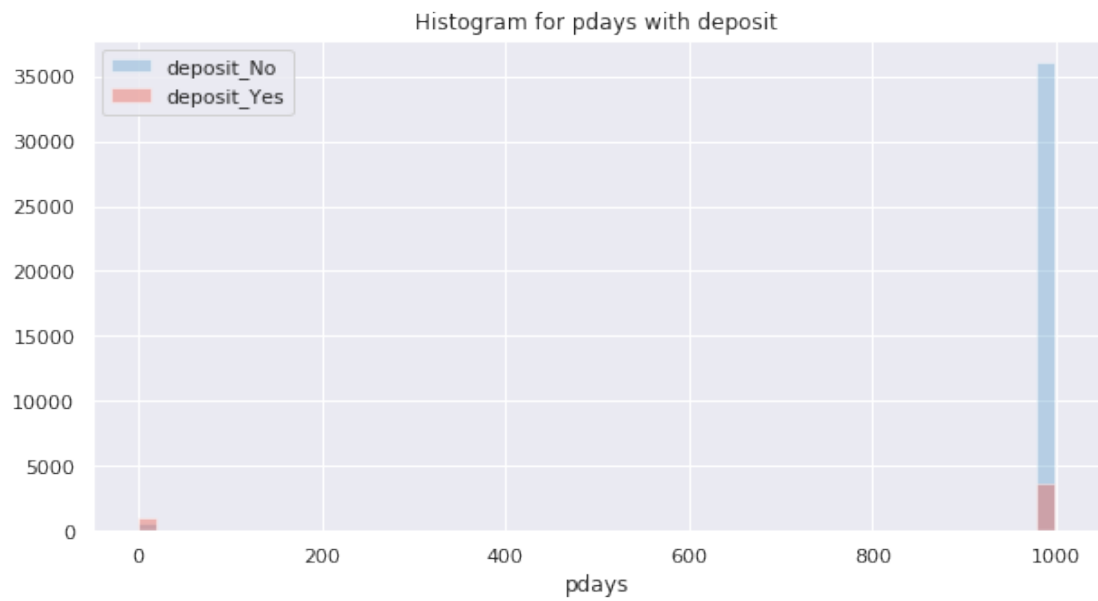
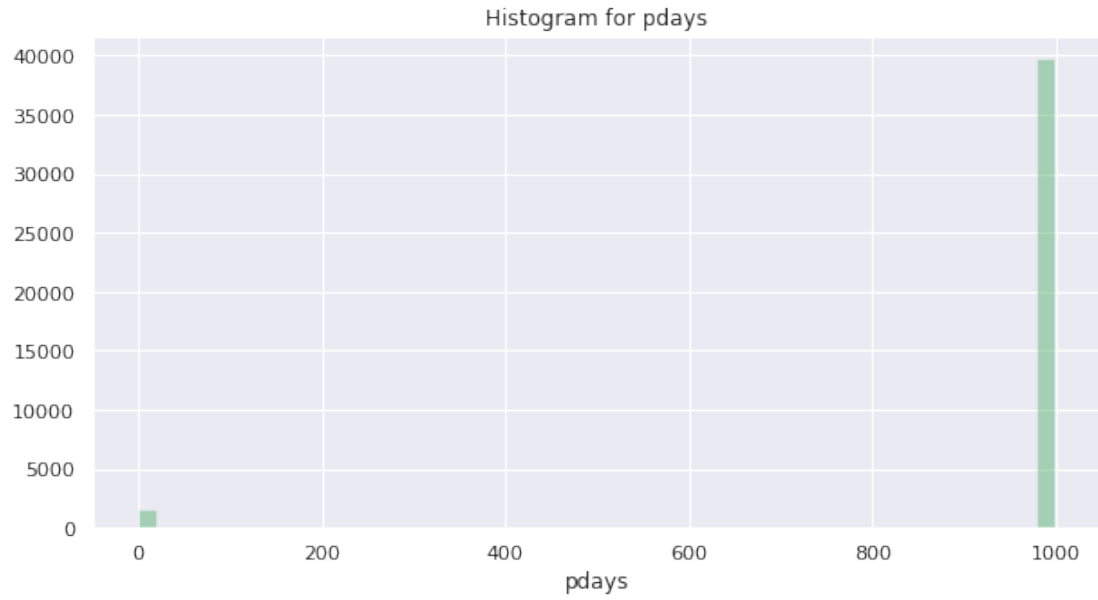
    def histogram_by_deposit(feature):
        plt.figure(figsize=(10, 5))
        plt.title("Histogram for {} with deposit".format(feature))
        ax0 = sns.distplot(dsdata[dsdata["y"]=="no"][feature], color=sns.color_palette("Blues_d")[4], kde=False, label="deposit_No")
        ax1 = sns.distplot(dsdata[dsdata["y"]=="yes"][feature], color=sns.color_palette("Reds_d")[4], kde=False, label="deposit_Yes")
        plt.legend()

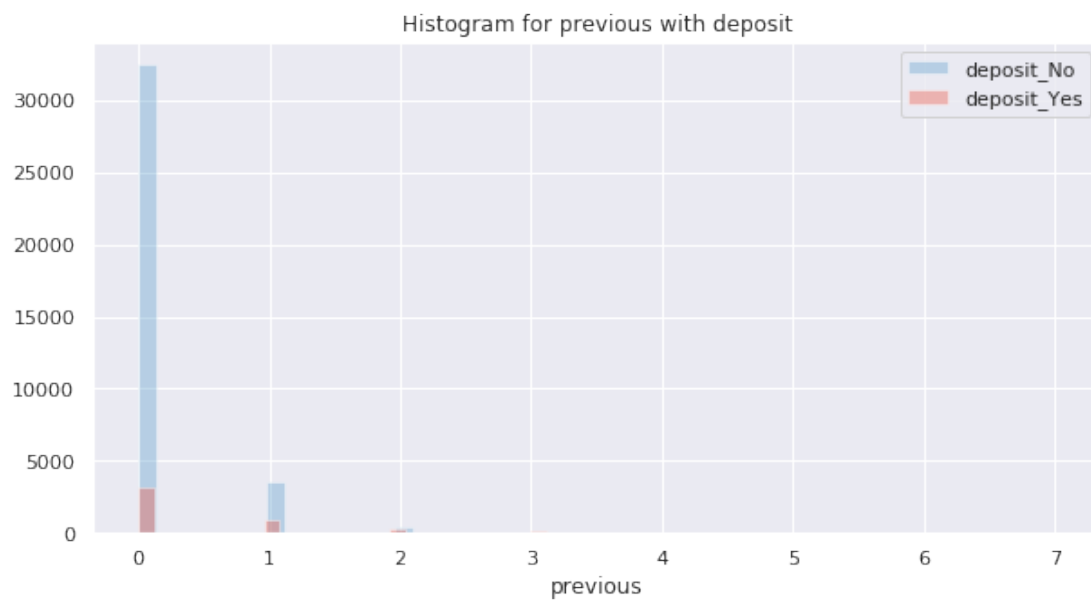
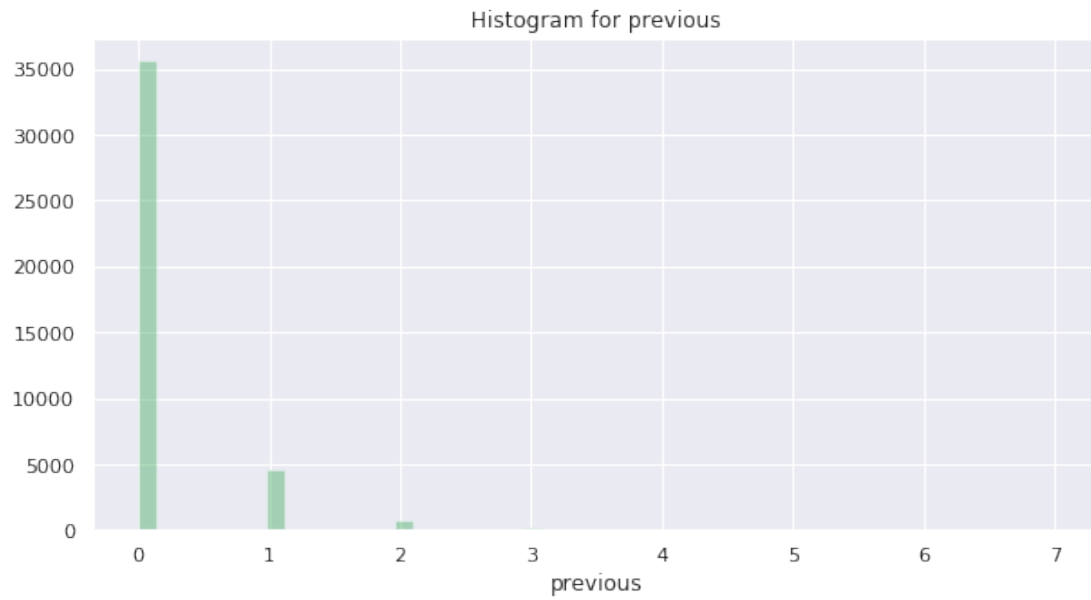
In [10]: # histogram for numerical variables
column = "job"
for column in dsdata.columns:
    if dsdata[column].dtype == "int64" or dsdata[column].dtype == "float64":
        histogram(column)
        histogram_by_deposit(column)
```

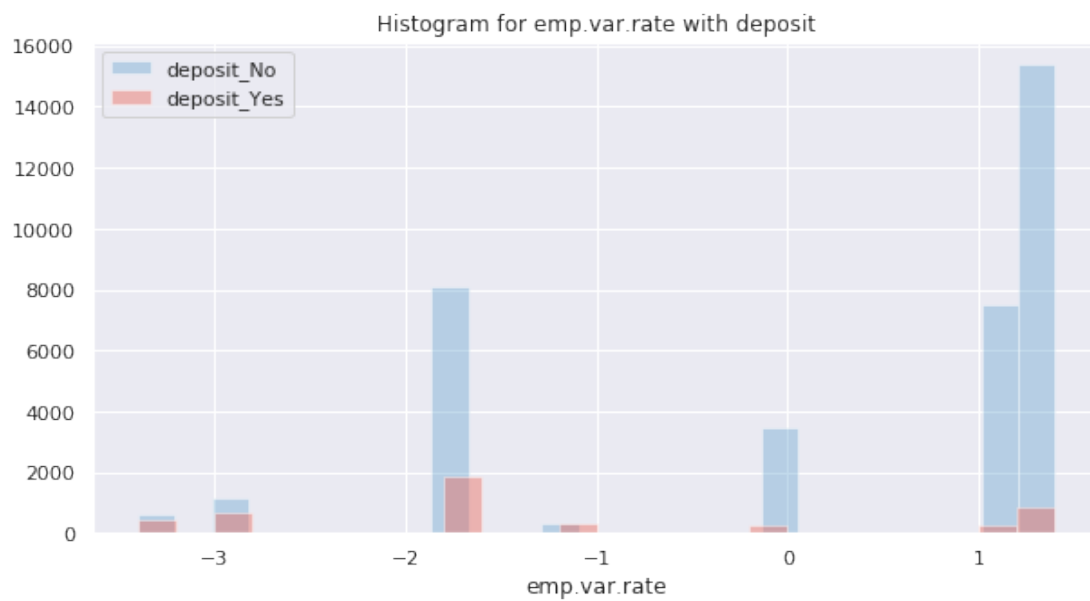
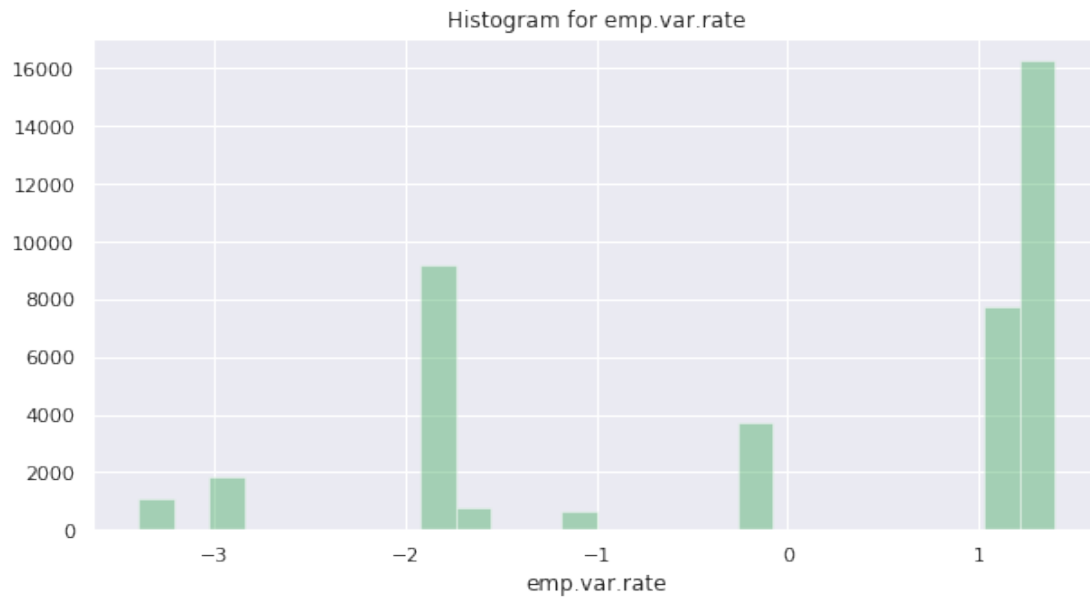






























III MISSING DATA HANDLING

In [11]: *#visualizing missing data*


```
msno.matrix(dsdata)  
plt.show()
```




In [12]: *#nullity correlation: how strongly the presence or absence of one variable affects the presence of another*
 msno.heatmap(dsdata)
 plt.show()
 missing = ['job', 'marital', 'education', 'default', 'housing', 'loan', 'poutcome']
 ds_missing = dsdata[missing]



In [13]: msno.bar(dsdata)
 plt.show()



```
In [14]: msno.dendrogram(dsdata)
plt.show()
```



```
In [15]: #ds_missing.head(25)
```

```
In [ ]:
```

```
In [16]: for col in list(ds_missing.columns):
print(ds_missing[col].value_counts(),'\n')
```

```
admin.          10422
blue-collar     9254
technician      6743
services        3969
management     2924
retired         1720
entrepreneur    1456
```

```
self-employed    1421
housemaid        1060
unemployed       1014
student          875
Name: job, dtype: int64
```

```
married    24928
single     11568
divorced   4612
Name: marital, dtype: int64
```

```
university.degree    12168
high.school           9515
basic.9y              6045
professional.course   5243
basic.4y              4176
basic.6y              2292
illiterate            18
Name: education, dtype: int64
```


```
no    32588
yes     3
Name: default, dtype: int64
```


```
yes    21576
no     18622
Name: housing, dtype: int64
```


```
no    33950
yes    6248
Name: loan, dtype: int64
```


```
failure    4252
success    1373
Name: poutcome, dtype: int64
```



```
In [17]: for col in list(ds_missing.columns):
          if ds_missing[col].dtype == "0":
              ds_missing[col].value_counts().plot(kind='bar')
              plt.title(col)
              plt.show()
```









```
In [18]: ds_missing.head(5)
```

```
Out[18]:
```

	job	marital	education	default	housing	loan	poutcome
0	housemaid	married	basic.4y	no	no	no	NaN
1	services	married	high.school	NaN	no	no	NaN
2	services	married	high.school	no	yes	no	NaN
3	admin.	married	basic.6y	no	no	no	NaN
4	services	married	high.school	no	no	yes	NaN

```
In [19]: missing2 = ['job', 'marital', 'education', 'housing', 'loan']
```

III Interpolation of categorical variables through empirical distributions


```
In [20]: dsdata2 = dsdata.copy()
dsdata2 = dsdata2.drop(columns=['poutcome','default']) #too many missing values, and
zero variance variable
#dsdata2.shape
```

```
In [21]: #fill missing data at random from discrete distribution corresponding to histogram
def dist_random_selection(col, num):
    arr = list(col.value_counts().index)
    prob = np.array(list(col.value_counts().values))
    p_norm = prob.sum()
    prob = prob/p_norm
    return np.random.choice(arr, num, replace=True, p=prob)

def fill_missing(data,missing):
    for col in list(data[missing].columns):
        count = len(data[col][data[col].isnull()])
        data[col][data[col].isnull()] = dist_random_selection(data[col],count)
    return data
```

```
In [22]: dsdata3 = fill_missing(dsdata2, missing2)
```

```
In [23]: msno.matrix(dsdata3)
plt.show()
```



IV FEATURE SELECTION

```
In [24]: dsdata3.shape
```

```
Out[24]: (41188, 19)
```

```
In [25]: # Data preparation
dsn = dsdata3.copy()
dsn.describe(include=['O'])
# Drop missing value
#dsn2 = dsn.copy().dropna() HELL NO!!!!!! XD,
# That drops about 10,000 columns...
```

```
Out[25]:
```

	job	marital	education	housing	loan	contact	month	\
count	41188	41188	41188	41188	41188	41188	41188	
unique	11	3	7	2	2	2	10	
top	admin.	married	university.degree	yes	no	cellular	may	
freq	10509	24977	12702	22121	34784	26144	13769	

	day_of_week	y
count	41188	41188
unique	5	2
top	thu	no
freq	8623	36548

```
In [26]: # Create dummy
def make_dummies(dsn):
    numvar = ['age', 'campaign', 'pdays', 'previous', 'emp.var.rate', 'cons.price.idx', 'cons.conf.idx', 'euribor3m', 'nr.employed']
    nonnumvar = ['job', 'marital', 'education', 'month', 'day_of_week'] #WE ONLY NEED THE CATEGORICALS, DONT INCLUDE BINARIES!!

    for c,var in enumerate(nonnumvar):
        dummy = pd.get_dummies(dsn[var],drop_first=True)
        dsn = dsn.drop(columns=[var])
        dsn = pd.concat([dsn, dummy], axis=1)
        #print(var)
        #display(dummy.head(5))
    return dsn

def make_numeric(dsn):
    for c,var in enumerate(['housing', 'loan', 'contact']):
        dsn[var] = dsn[var].astype("category").cat.codes

    return dsn

dsn2 = make_dummies(dsn) #make dummies out of categoricals
dsn2 = make_numeric(dsn2) #make binaries out of yes/no
dsn2.info()
#dsn2.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41188 entries, 0 to 41187
Data columns (total 45 columns):
age                41188 non-null int64
housing            41188 non-null int8
loan               41188 non-null int8
```

```

contact          41188 non-null int8
duration         41188 non-null int64
campaign         41188 non-null int64
pdays          41188 non-null int64
previous         41188 non-null int64
emp.var.rate    41188 non-null float64
cons.price.idx  41188 non-null float64
cons.conf.idx   41188 non-null float64
euribor3m       41188 non-null float64
nr.employed     41188 non-null float64
y               41188 non-null object
blue-collar     41188 non-null uint8
entrepreneur    41188 non-null uint8
housemaid       41188 non-null uint8
management      41188 non-null uint8
retired         41188 non-null uint8
self-employed   41188 non-null uint8
services        41188 non-null uint8
student         41188 non-null uint8
technician      41188 non-null uint8
unemployed      41188 non-null uint8
married         41188 non-null uint8
single          41188 non-null uint8
basic.6y        41188 non-null uint8
basic.9y        41188 non-null uint8
high.school     41188 non-null uint8
illiterate      41188 non-null uint8
professional.course 41188 non-null uint8
university.degree 41188 non-null uint8
aug             41188 non-null uint8
dec             41188 non-null uint8
jul             41188 non-null uint8
jun             41188 non-null uint8
mar             41188 non-null uint8
may             41188 non-null uint8
nov             41188 non-null uint8
oct             41188 non-null uint8
sep             41188 non-null uint8
mon             41188 non-null uint8
thu             41188 non-null uint8
tue             41188 non-null uint8
wed             41188 non-null uint8
dtypes: float64(5), int64(5), int8(3), object(1), uint8(31)
memory usage: 4.8+ MB

```

```

In [27]: # Drop priori
         dsn3 = dsn2.drop(columns=['duration']) #remove output and duration, which should not be
         known a priori

```

```

In [28]: # Data Normalization

         # separate the data from the target attributes
         X = dsn3.drop(columns=['y'])
         #Y = pd.get_dummies(dsn3['y'], drop_first=True, dummy_na=True)

         # normalize the data attributes

```



```
normalized_X = preprocessing.normalize(X)
```

```
In [29]: # MAKE NORMALIZED DF
X_n = pd.DataFrame(normalized_X)
X_n.columns = X.columns
```


```
In [30]: # Check the correlation

# NORMALIZE THE DATA !
corr = X_n.corr()
# Generate a mask for the upper triangle
sns.set(style="white")
mask = np.zeros_like(corr, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True

# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(11, 9))

# Generate a custom diverging colormap
cmap = sns.diverging_palette(220, 10, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=.3, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .76})
f.show()
```



IV Feature Selection

```
In [31]: def print_VIF(X):
          colnames = list(X.columns)
          for i in range(X.shape[1]):
              print(oi.variance_inflation_factor(X.values, i), colnames[i])

          print("Variance Inflation Factors:")
          print_VIF(X)
```

```
Variance Inflation Factors:
25.699421717728164 age
2.187347297019118 housing
1.1884217905667915 loan
5.168655621475389 contact
1.9403265754912258 campaign
45.31697288715908 pdays
2.0804063875271255 previous
93.67968903223512 emp.var.rate
62381.33728205274 cons.price.idx
```

```

389.93355428533994 cons.conf.idx
785.1827162631885 euribor3m
80095.1601608278 nr.employed
2.998359623841996 blue-collar
1.1801111052746007 entrepreneur
1.208239433519948 housemaid
1.3392985940698983 management
1.5422587461133799 retired
1.1596803634308726 self-employed
1.5653551621374622 services
1.1949262756217893 student
2.077822521405796 technician
1.124826196191629 unemployed
6.547908953431982 married
4.031673284036735 single
1.598731702634765 basic.6y
2.6330516709015193 basic.9y
4.342608264937074 high.school
1.0051487332377531 illiterate
3.007594830576488 professional.course
5.7058176510336525 university.degree
8.131916021936958 aug
1.1411179834728389 dec
5.437300515663982 jul
3.78227309374364 jun
1.2543889589126438 mar
7.5898279431870295 may
3.9328083828264737 nov
1.5991916740322096 oct
1.518849245189762 sep
2.0996206920666816 mon
2.118840553704938 thu
2.056199982646952 tue
2.0552587590547184 wed


```

1. intermediate regression on economic variables for VIF

```

In [32]: economic = ['emp.var.rate', 'cons.price.idx', 'cons.conf.idx', 'euribor3m']
         numeric = sm.OLS(X['nr.employed'].values, X[economic].values).fit()
         sns.distplot(numeric.resid)
         plt.title("Residuals")
         plt.show()
         display(numeric.summary())

```



```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

OLS Regression Results

```
=====
Dep. Variable:          y      R-squared:                1.000
Model:                  OLS    Adj. R-squared:            1.000
Method:                 Least Squares    F-statistic:        2.623e+08
Date:                   Fri, 10 May 2019    Prob (F-statistic):    0.00
Time:                   14:05:52    Log-Likelihood:       -2.0167e+05
No. Observations:       41188    AIC:                  4.034e+05
Df Residuals:           41184    BIC:                  4.034e+05
Df Model:                4
Covariance Type:        nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
x1	-63.3977	0.445	-142.626	0.000	-64.269	-62.526
x2	50.2830	0.027	1890.625	0.000	50.231	50.335
x3	-3.7858	0.038	-99.832	0.000	-3.860	-3.711
x4	86.5789	0.416	208.289	0.000	85.764	87.394

```
=====
Omnibus:                 3472.453    Durbin-Watson:           0.002
Prob(Omnibus):            0.000    Jarque-Bera (JB):        9896.607
Skew:                     -0.464    Prob(JB):                 0.00
Kurtosis:                  5.215    Cond. No.                  387.
=====
```

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
"""
```

```
In [33]: dsn4 = X.copy()
        dsn4 = X.drop(columns=['nr.employed', 'euribor3m', 'cons.price.idx', 'cons.conf.idx'])

        print("Variance Inflation Factors:")
        print_VIF(dsn4)
```

```
Variance Inflation Factors:
18.795212167761804 age
2.1680344085946066 housing
1.186587839574958 loan
3.764577280072294 contact
1.9215346653769463 campaign
27.27979619302259 pdays
1.7329058424985795 previous
2.5131426179783563 emp.var.rate
2.826150727152902 blue-collar
1.1735843768089205 entrepreneur
1.192256422293572 housemaid
1.3361058353560156 management
1.5358166880014659 retired
1.1549681725521808 self-employed
1.541617342271284 services
1.1675162611834708 student
2.046333858599015 technician
1.1167136673732803 unemployed
6.030824907822872 married
3.538584478475322 single
1.5287117971689477 basic.6y
2.4156871535745448 basic.9y
3.7703535660523375 high.school
1.0046374774886204 illiterate
2.7470776917530637 professional.course
4.867456484498119 university.degree
3.91727535314813 aug
1.0774840887974682 dec
4.41102740372729 jul
3.358396917545775 jun
1.1875430357287888 mar
6.319548041788819 may
2.5120859115787244 nov
1.273633542385804 oct
1.2196007096354973 sep
2.051357368207603 mon
2.058402558227157 thu
2.0140216689556043 tue
2.0109137649930435 wed
```

```
In [34]: dsn4 = dsn4.drop(columns=['pdays'])
        dsn4 = dsn4.drop(columns=['mon', 'thu', 'tue', 'wed'])
        dsn4 = dsn4.drop(columns=['single'])
        #dsn4 = dsn4.drop(columns=['campaign'])

        # NORMALIZE THE DATA !
        normalized_X = preprocessing.normalize(dsn4)

        X_n = pd.DataFrame(normalized_X)
```

```
X_n.columns = dsn4.columns
```


```
#plt.figure(figsize=(12,8))
corr = X_n.corr()
```

```
# Generate a mask for the upper triangle
sns.set(style="white")
mask = np.zeros_like(corr, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True
```

```
# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(11, 9))
```

```
# Generate a custom diverging colormap
cmap = sns.diverging_palette(220, 10, as_cmap=True)
```

```
# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=.3, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .75})
plt.show()
```



IV Outlier Detection

1. PCA

```
In [35]: from sklearn.decomposition import PCA
         from sklearn.preprocessing import StandardScaler

In [ ]:

In [36]: X = dsn4#dsn4.drop(columns=['y'])

         #sns.pairplot(X)
         # Plot the data
         #fig = plt.figure(figsize=(12,8))
         #with plt.style.context(('ggplot')):
         #    plt.plot( X.T)
         #    plt.show()

In [ ]:

In [37]: pcaA = PCA()
         pcaX = pcaA.fit_transform(X)#.fit_transform(StandardScaler().fit_transform(X))
         # PCA & score
         print(pcaA.explained_variance_ratio_[0:5])

[0.89099929 0.06367829 0.02046383 0.00326929 0.00232139]

In [38]: # Compute the euclidean distance ( 3 PC )
         euclidean = np.zeros(X.shape[0])
         for i in range(3):
             euclidean += (pcaX[:,i] - np.mean(pcaX[:,3]))**2/np.var(pcaX[:,3])

         #colors = [plt.cm.jet(float(i)/max(euclidean)) for i in euclidean]

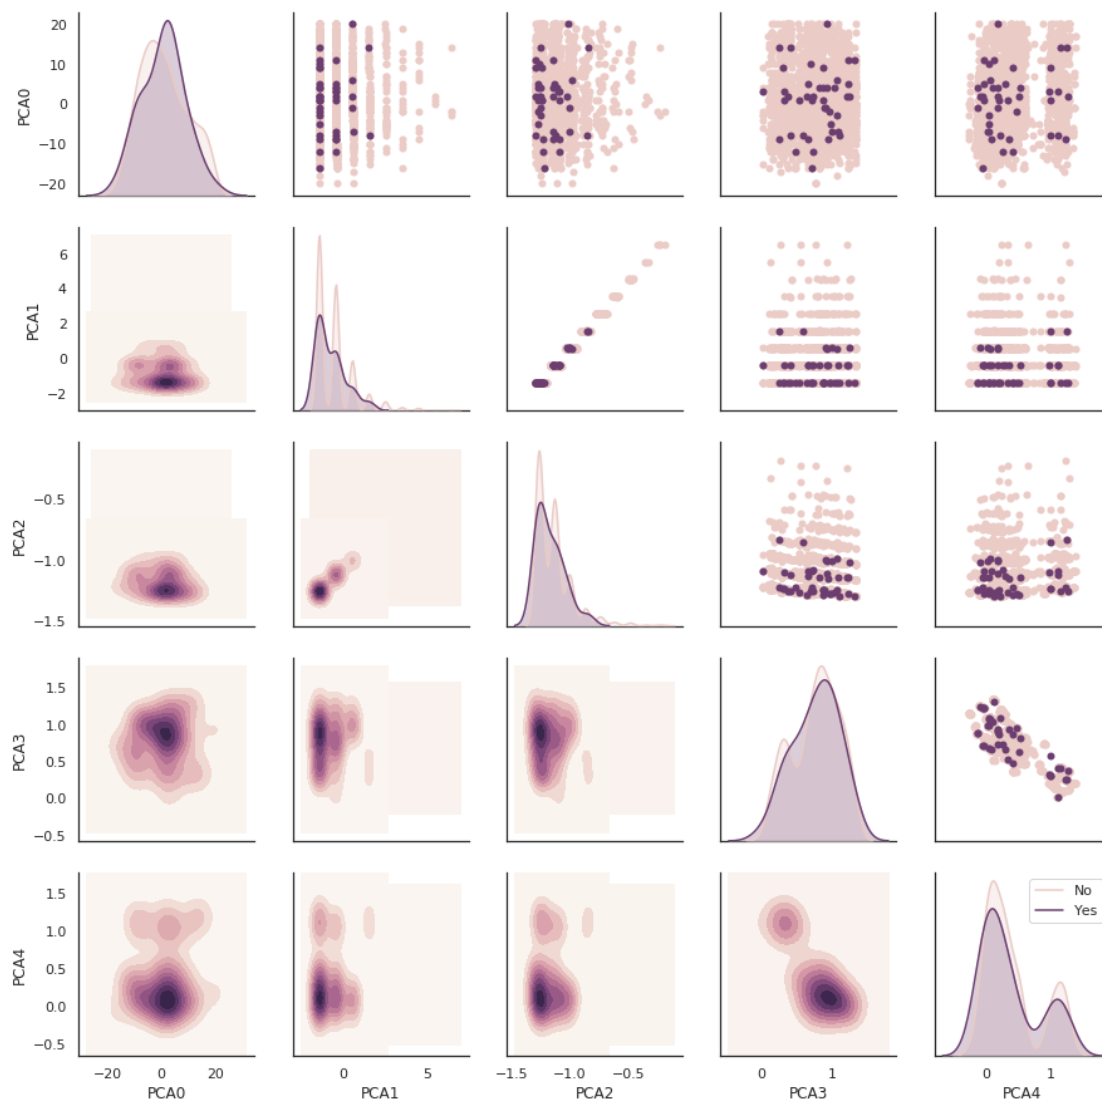
In [39]: X_pca = pd.DataFrame(pcaX, columns=['PCA%i' % i for i in range(X.shape[1])],
         index=X.index)

In [40]: colors = ["blue", "blue"]
         pal = sns.xkcd_palette(colors)

         pal = [sns.cubehelix_palette(light=1)[1], sns.cubehelix_palette(light=1)[4]]

In [41]: sns.set_style('white')
         cmap = sns.cubehelix_palette(light=1, as_cmap=True)

         df = X_pca.copy()
         df['y'] = dsdata['y']
         df = df[['PCA0', 'PCA1', 'PCA2', 'PCA3', 'PCA4', 'y']]
         df = df.iloc[:2000]
         g = sns.PairGrid(df, diag_sharey=False, hue='y', palette=pal)
         g.map_lower(sns.kdeplot, cmap=cmap, shade=True)
         g.map_upper(sns.scatterplot, linewidth=0)
         g.map_diag(sns.kdeplot, shade=True)
         plt.legend(['No', 'Yes'])
         plt.show()
```



```
In [42]: '''
          colors = [plt.cm.jet(float(i)/max(euclidean)) for i in euclidean]
          fig = plt.figure(figsize=(8,6))
          with plt.style.context('ggplot'):
              plt.scatter(pcaX[:, 0], pcaX[:, 1], c=colors, edgecolors='k')
              plt.xlabel('PC1')
              plt.ylabel('PC2')
              plt.title('Score Plot')
          plt.show()
          '''
          a = 0
```

2. Z-score

```
In [43]: #NORMALIZATION????
```



```
In [44]: from scipy import stats
```

```
        dsn5 = dsn4.copy()
```


```
        zX = np.abs(stats.zscore(X))
```

```
In [45]: X.shape
```

```
Out[45]: (41188, 33)
```

```
In [46]: def histogram(variable):
        plt.figure(figsize=(10, 5))
        plt.title("Histogram")
        ax = plt.hist(zX)
```

```
        histogram(zX)
```



```
In [47]: dsn5 = dsn5[(zX < 10).all(axis=1)]
        dsn5 = dsn5.drop(columns = ['dec', 'illiterate'])
        dsn5.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 40949 entries, 0 to 41187
Data columns (total 31 columns):
age                40949 non-null int64
housing            40949 non-null int8
loan               40949 non-null int8
contact            40949 non-null int8
campaign           40949 non-null int64
previous           40949 non-null int64
emp.var.rate       40949 non-null float64
blue-collar        40949 non-null uint8
entrepreneur       40949 non-null uint8
housemaid          40949 non-null uint8
management         40949 non-null uint8
retired            40949 non-null uint8
```

```

self-employed      40949 non-null uint8
services           40949 non-null uint8
student            40949 non-null uint8
technician         40949 non-null uint8
unemployed         40949 non-null uint8
married            40949 non-null uint8
basic.6y           40949 non-null uint8
basic.9y           40949 non-null uint8
high.school        40949 non-null uint8
professional.course 40949 non-null uint8
university.degree  40949 non-null uint8
aug                40949 non-null uint8
jul                40949 non-null uint8
jun                40949 non-null uint8
mar                40949 non-null uint8
may                40949 non-null uint8
nov                40949 non-null uint8
oct                40949 non-null uint8
sep                40949 non-null uint8
dtypes: float64(1), int64(3), int8(3), uint8(24)
memory usage: 2.6 MB

```

```

In [48]: X = dsn5#.drop(columns=['y'])
         Y = dsdata['y'][X.index].astype("category").cat.codes
         Y.value_counts()

```

```

Out[48]: 0      36405
         1       4544
         dtype: int64

```

V PREDICTION

```

In [49]: sns.set_style('whitegrid')

```

```

In [ ]:

```


```

In [50]: X_ = sm.add_constant(X)
         model = sm.OLS(Y,X_).fit()

         print("Distribution of OLS residuals")
         sns.distplot(model.resid)
         plt.show()
         display(model.summary())

```

Distribution of OLS residuals



```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

OLS Regression Results

```
=====
Dep. Variable:          y      R-squared:                0.144
Model:                  OLS    Adj. R-squared:            0.144
Method:                 Least Squares    F-statistic:          222.7
Date:                   Fri, 10 May 2019    Prob (F-statistic):    0.00
Time:                   14:06:10    Log-Likelihood:        -7489.8
No. Observations:       40949    AIC:                   1.504e+04
Df Residuals:           40917    BIC:                   1.532e+04
Df Model:               31
Covariance Type:        nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	0.0891	0.012	7.683	0.000	0.066	0.112
age	0.0003	0.000	1.696	0.090	-4.42e-05	0.001
housing	-0.0049	0.003	-1.701	0.089	-0.011	0.001
loan	-0.0021	0.004	-0.520	0.603	-0.010	0.006
contact	0.0288	0.005	6.192	0.000	0.020	0.038
campaign	-0.0035	0.001	-6.257	0.000	-0.005	-0.002
previous	0.0686	0.003	20.832	0.000	0.062	0.075
emp.var.rate	-0.0529	0.001	-36.062	0.000	-0.056	-0.050
blue-collar	-0.0218	0.005	-4.178	0.000	-0.032	-0.012
entrepreneur	-0.0210	0.008	-2.546	0.011	-0.037	-0.005
housemaid	-0.0076	0.010	-0.776	0.438	-0.027	0.012
management	-0.0127	0.006	-2.046	0.041	-0.025	-0.001
retired	0.0438	0.009	5.026	0.000	0.027	0.061
self-employed	-0.0155	0.008	-1.865	0.062	-0.032	0.001
services	-0.0178	0.006	-3.094	0.002	-0.029	-0.007

student	0.0771	0.011	7.178	0.000	0.056	0.098
technician	-0.0070	0.005	-1.366	0.172	-0.017	0.003
unemployed	0.0019	0.010	0.199	0.842	-0.017	0.021
married	-0.0033	0.003	-1.062	0.288	-0.010	0.003
basic.6y	0.0013	0.008	0.167	0.868	-0.014	0.016
basic.9y	-0.0083	0.006	-1.397	0.162	-0.020	0.003
high.school	-0.0042	0.006	-0.686	0.493	-0.016	0.008
professional.course	0.0013	0.007	0.183	0.855	-0.012	0.015
university.degree	0.0090	0.006	1.456	0.145	-0.003	0.021
aug	0.0495	0.008	6.375	0.000	0.034	0.065
jul	0.0677	0.008	8.752	0.000	0.053	0.083
jun	0.0342	0.008	4.465	0.000	0.019	0.049
mar	0.2778	0.014	20.256	0.000	0.251	0.305
may	-0.0495	0.007	-7.536	0.000	-0.062	-0.037
nov	-0.0273	0.008	-3.622	0.000	-0.042	-0.013
oct	0.1579	0.012	12.753	0.000	0.134	0.182
sep	0.1714	0.014	12.625	0.000	0.145	0.198

```
=====
Omnibus:                16009.730    Durbin-Watson:                1.812
Prob(Omnibus):          0.000    Jarque-Bera (JB):            54082.712
Skew:                   2.043    Prob(JB):                     0.00
Kurtosis:               6.873    Cond. No.                     570.
=====
```

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
"""
```



```
In [51]: alpha = 0.05
         a = model.pvalues < alpha

         X2 = X_[X_.columns[a]]
         X2 = sm.add_constant(X2)
         print("Not Statistically significant regressors are:")
         print(list(X_.columns[~a]))
```

Not Statistically significant regressors are:

```
['age', 'housing', 'loan', 'housemaid', 'self-employed', 'technician', 'unemployed',
'married', 'basic.6y', 'basic.9y', 'high.school', 'professional.course',
'university.degree']
```

```
In [52]: model2 = sm.OLS(Y,X2).fit(cov_type='HCO')
         sm.qqplot(model2.resid, sc.norm, fit=True, line='45')
         plt.show()
         sns.distplot(model2.resid)
         plt.title('Residuals')
         plt.show()
         display(model2.summary())
         print("Variance Inflation Factors:")
         print_VIF(X2)
```



```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

OLS Regression Results

=====

```

Dep. Variable:          y      R-squared:          0.144
Model:                OLS      Adj. R-squared:       0.143
Method:             Least Squares  F-statistic:       188.4
Date:                Fri, 10 May 2019  Prob (F-statistic):    0.00
Time:                14:06:11   Log-Likelihood:     -7503.5
No. Observations:    40949      AIC:               1.505e+04
Df Residuals:        40930      BIC:               1.521e+04
Df Model:            18
Covariance Type:      HCO

```

```

=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
const          0.0936      0.009     10.393      0.000      0.076      0.111
contact        0.0294      0.006      5.296      0.000      0.019      0.040
campaign       -0.0035      0.000     -8.867      0.000     -0.004     -0.003
previous       0.0688      0.005     13.464      0.000      0.059      0.079
emp.var.rate   -0.0531      0.002    -25.669      0.000     -0.057     -0.049
blue-collar    -0.0238      0.003     -7.031      0.000     -0.030     -0.017
entrepreneur   -0.0173      0.007     -2.348      0.019     -0.032     -0.003
management     -0.0055      0.006     -0.945      0.345     -0.017      0.006
retired        0.0509      0.010      5.147      0.000      0.032      0.070
services       -0.0197      0.005     -4.166      0.000     -0.029     -0.010
student        0.0751      0.015      4.980      0.000      0.046      0.105
aug            0.0515      0.011      4.799      0.000      0.030      0.072
jul            0.0677      0.010      6.508      0.000      0.047      0.088
jun            0.0342      0.010      3.406      0.001      0.015      0.054
mar            0.2794      0.023     12.308      0.000      0.235      0.324
may           -0.0504      0.009     -5.821      0.000     -0.067     -0.033
nov           -0.0268      0.010     -2.808      0.005     -0.046     -0.008
oct            0.1589      0.021      7.690      0.000      0.118      0.199
sep            0.1730      0.023      7.647      0.000      0.129      0.217
=====
Omnibus:                16022.479   Durbin-Watson:           1.810
Prob(Omnibus):           0.000   Jarque-Bera (JB):       54146.902
Skew:                    2.045   Prob(JB):                0.00
Kurtosis:                 6.874   Cond. No.                50.7
=====

```

Warnings:

```

[1] Standard Errors are heteroscedasticity robust (HCO)
"""

```

Variance Inflation Factors:

```

21.134835195011743 const
2.4230593529335542 contact
1.035602184712304 campaign
1.2461663846573323 previous
2.525902975234939 emp.var.rate
1.176311232198173 blue-collar
1.0429912415669762 entrepreneur
1.0693348771898068 management
1.0610369988444968 retired
1.0977835821445638 services
1.0476733948390058 student
3.718192758459077 aug
4.160433242738336 jul

```

```

3.2062143149574265 jun
1.1978096592691139 mar
4.6599163715584835 may
2.4696254452997293 nov
1.2793457906752723 oct
1.2250455192107061 sep

```

```
In [53]: X_train, X_test, Y_train, Y_test = train_test_split(X2.values, Y.values, test_size=0.25)
```

```

reg = LinearRegression()
reg = reg.fit(X_train,Y_train)

```

```
y_hat = reg.predict(X_test)
```

```

print('Test accuracy:',np.round(reg.score(np.round(X_test), Y_test),3), ', MSE Loss
is:', mean_squared_error(Y_test,y_hat))

```

```
Test accuracy: 0.145 , MSE Loss is: 0.08108139598792455
```

```
In [ ]:
```

VI CLASSIFICATION

```
In [54]: def plot_hist2(df,df2,df_col):
```

```
    df = df.dropna()
```

```
    df2 = df2.dropna()
```

```
    for d in df_col:
```

```
        print("Empirical Distribution of Variable "+d)
```

```
        fig, axes = plt.subplots(1,2,figsize=(15,9))
```

```
        sns.distplot(df[d],ax=axes[0])
```

```
        sns.distplot(df2[d],ax=axes[1])
```

```
        axes[0].set_ylabel("Probability")
```

```
        axes[1].set_ylabel("Probability")
```

```
        plt.suptitle("Empirical Probability Distribution of Numerical Variable "+d)
```

```
        plt.show()
```

```
def get_num_cols(df):
```

```
    idx = df.select_dtypes(exclude='object').columns.values
```

```
    dF = df[idx].dropna() #remove NaNs or else it cant plot
```

```
    return dF.columns
```

```
def calculate_metrics(y_test,y_hat):
```

```
    c = confusion_matrix(y_test, y_hat)
```

```
    print("Confusion matrix is:")
```

```
    print(c)
```

```
    print("We have",c[0][0]+c[1][1],"correct observations and",c[0][1]+c[1][0],
"misclassifications.")
```

```
    print(classification_report(y_test, y_hat))
```

```
    plt.figure(figsize=(6,6))
```

```
    sns.heatmap(c,cmap="YlGnBu",annot=True,fmt='g')
```

```
    plt.show()
```

```
def plot_ROC(y_test,X_test,classifier):
```

```

roc = roc_auc_score(y_test, classifier.predict(X_test))
fpr, tpr, _ = roc_curve(y_test, classifier.predict_proba(X_test)[: ,1])
plt.figure(figsize=(12,12))
plt.plot(fpr, tpr, label='Classifier area =' +str(np.round(roc,2)))
plt.plot([0, 1], [0, 1], '--')

plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('ROC curve')
plt.legend()

plt.xlim([-0.01, 1.0])
plt.ylim([0.0, 1.01])
plt.show()

def plot_ROC2(L_Y,L_X,L_YHAT,L_PROB,L_NAME):

    plt.figure(figsize=(12,12))
    for i in range(len(L_Y)):
        roc = roc_auc_score(L_Y[i], L_YHAT[i])
        fpr, tpr, _ = roc_curve(L_Y[i], L_PROB[i] [:,1])
        plt.plot(fpr, tpr, label=L_NAME[i]+' , area =' +str(np.round(roc,2)))
    plt.plot([0, 1], [0, 1], '--')

    plt.xlabel('FPR')
    plt.ylabel('TPR')
    plt.title('ROC curve')
    plt.legend()

    plt.xlim([-0.01, 1.0])
    plt.ylim([0.0, 1.01])
    plt.show()

def MAE(y_test,y_hat):
    return np.abs(y_test-y_hat).sum()#/y_test.shape[0]

```

VI Logistic Regression with PCA

```

In [55]: pcaXmodel = PCA(n_components=5, whiten=True)
pc = pcaXmodel.fit_transform(X) #.fit_transform(StandardScaler().fit_transform(X))
print(pcaXmodel.explained_variance_ratio_[:5])

proj = pcaXmodel.inverse_transform(pc)
a = pd.DataFrame(proj)[list((np.ones((31,1))-1).cumsum().ravel())]

X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.3,
random_state=0)
lr = LogisticRegression().fit(X_train,y_train)

y_hat = lr.predict(X_test)
print('Test accuracy:',np.round(lr.score(X_test, y_test),3), ', Cross Entropy Loss is:',
log_loss(y_test,y_hat))

```



```
print(MAE(y_test,y_hat))

#pcaXmodel =
PCA(n_components=5).fit_transform(X)#.fit_transform(StandardScaler().fit_transform(X))
```

```
[0.89682587 0.05748246 0.0204958 0.00333079 0.00236176]
Test accuracy: 0.887 , Cross Entropy Loss is: 3.905127582100686
1389
```


```
In [56]: calculate_metrics(y_test,y_hat)
         plot_ROC(y_test,X_test,lr)
```


Confusion matrix is:

```
[[10737  165]
 [ 1224  159]]
```

We have 10896 correct observations and 1389 misclassifications.

	precision	recall	f1-score	support
0	0.90	0.98	0.94	10902
1	0.49	0.11	0.19	1383
micro avg	0.89	0.89	0.89	12285
macro avg	0.69	0.55	0.56	12285
weighted avg	0.85	0.89	0.85	12285





```
In [57]: a.shape
```

```
Out[57]: (40949, 31)
```

```
In [58]: logit = sm.Logit(Y, X).fit()
display(logit.summary())
```

```
Optimization terminated successfully.
Current function value: 0.293562
Iterations 7
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

Logit Regression Results

```
=====
Dep. Variable:          y    No. Observations:      40949
Model:                Logit    Df Residuals:      40918
Method:                MLE     Df Model:         30
```

```

Date:          Fri, 10 May 2019   Pseudo R-squ.:      0.1577
Time:          14:06:12          Log-Likelihood:    -12021.
converged:     True              LL-Null:              -14272.
                                   LLR p-value:          0.000

```

	coef	std err	z	P> z	[0.025	0.975]
age	-0.0208	0.001	-13.966	0.000	-0.024	-0.018
housing	-0.1384	0.034	-4.108	0.000	-0.204	-0.072
loan	-0.0677	0.048	-1.424	0.154	-0.161	0.025
contact	-0.2694	0.051	-5.239	0.000	-0.370	-0.169
campaign	-0.0805	0.010	-8.272	0.000	-0.100	-0.061
previous	0.3956	0.026	15.024	0.000	0.344	0.447
emp.var.rate	-0.4211	0.014	-31.187	0.000	-0.448	-0.395
blue-collar	-0.6715	0.060	-11.141	0.000	-0.790	-0.553
entrepreneur	-0.3916	0.105	-3.731	0.000	-0.597	-0.186
housemaid	-0.3927	0.121	-3.249	0.001	-0.630	-0.156
management	-0.2096	0.072	-2.903	0.004	-0.351	-0.068
retired	0.4508	0.089	5.043	0.000	0.276	0.626
self-employed	-0.3553	0.099	-3.579	0.000	-0.550	-0.161
services	-0.4390	0.071	-6.214	0.000	-0.577	-0.301
student	-0.0562	0.089	-0.630	0.529	-0.231	0.119
technician	-0.2691	0.058	-4.633	0.000	-0.383	-0.155
unemployed	-0.2013	0.106	-1.902	0.057	-0.409	0.006
married	-0.0648	0.037	-1.746	0.081	-0.138	0.008
basic.6y	-0.4636	0.091	-5.101	0.000	-0.642	-0.285
basic.9y	-0.6431	0.067	-9.652	0.000	-0.774	-0.512
high.school	-0.7244	0.059	-12.370	0.000	-0.839	-0.610
professional.course	-0.5780	0.071	-8.130	0.000	-0.717	-0.439
university.degree	-0.6507	0.056	-11.619	0.000	-0.760	-0.541
aug	-0.2108	0.065	-3.234	0.001	-0.338	-0.083
jul	0.0587	0.066	0.890	0.373	-0.071	0.188
jun	-0.0449	0.069	-0.655	0.512	-0.179	0.089
mar	0.9619	0.098	9.778	0.000	0.769	1.155
may	-0.9716	0.056	-17.424	0.000	-1.081	-0.862
nov	-0.6523	0.070	-9.367	0.000	-0.789	-0.516
oct	0.2948	0.093	3.183	0.001	0.113	0.476
sep	0.4243	0.100	4.253	0.000	0.229	0.620

```

"""

```

```

In [59]: sns.set_style('whitegrid')
         alpha = 0.05
         a = logit.pvalues < alpha

         X3 = X[X.columns[a]]
         print("Not Statistically significant regressors are:")
         print(list(X.columns[~a]))


Not Statistically significant regressors are:
['loan', 'student', 'unemployed', 'married', 'jul', 'jun']


In [60]: model2 = sm.Logit(Y.values,X3.values).fit()
         sm.qqplot(model2.resid_dev, sc.norm, fit=True, line='45')
         plt.grid()

```

```
plt.show()
sns.distplot(model2.resid_dev)
plt.title('Residuals')
plt.grid()
plt.show()
display(model2.summary())
print("Variance Inflation Factors:")
print_VIF(X3)
```

Optimization terminated successfully.
Current function value: 0.293698
Iterations 7





```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

Logit Regression Results

```
=====
Dep. Variable:          y      No. Observations:      40949
Model:                Logit      Df Residuals:      40924
Method:                MLE       Df Model:         24
Date:                Fri, 10 May 2019      Pseudo R-squ.:      0.1573
Time:                14:06:12      Log-Likelihood:     -12027.
converged:              True      LL-Null:          -14272.
                                LLR p-value:           0.000
=====
```

	coef	std err	z	P> z	[0.025	0.975]
x1	-0.0219	0.001	-16.668	0.000	-0.024	-0.019
x2	-0.1418	0.034	-4.226	0.000	-0.208	-0.076
x3	-0.3036	0.047	-6.481	0.000	-0.395	-0.212
x4	-0.0805	0.010	-8.349	0.000	-0.099	-0.062
x5	0.3958	0.026	15.082	0.000	0.344	0.447
x6	-0.4135	0.012	-34.306	0.000	-0.437	-0.390
x7	-0.6559	0.058	-11.275	0.000	-0.770	-0.542
x8	-0.3819	0.104	-3.668	0.000	-0.586	-0.178
x9	-0.3692	0.120	-3.084	0.002	-0.604	-0.135
x10	-0.2024	0.071	-2.842	0.004	-0.342	-0.063
x11	0.4948	0.088	5.640	0.000	0.323	0.667
x12	-0.3397	0.099	-3.446	0.001	-0.533	-0.146
x13	-0.4211	0.069	-6.071	0.000	-0.557	-0.285
x14	-0.2481	0.057	-4.384	0.000	-0.359	-0.137
x15	-0.4730	0.090	-5.268	0.000	-0.649	-0.297
x16	-0.6526	0.065	-10.086	0.000	-0.779	-0.526
x17	-0.7189	0.055	-13.006	0.000	-0.827	-0.611

x18	-0.5823	0.069	-8.402	0.000	-0.718	-0.446
x19	-0.6397	0.053	-12.089	0.000	-0.743	-0.536
x20	-0.2317	0.054	-4.328	0.000	-0.337	-0.127
x21	0.9615	0.093	10.315	0.000	0.779	1.144
x22	-0.9750	0.045	-21.635	0.000	-1.063	-0.887
x23	-0.6670	0.061	-10.973	0.000	-0.786	-0.548
x24	0.2965	0.088	3.378	0.001	0.124	0.469
x25	0.4162	0.095	4.400	0.000	0.231	0.602

```
=====
"""
```

Variance Inflation Factors:

```
10.59285135212034 age
2.1189581199142262 housing
2.5784705320747205 contact
1.9652236412758264 campaign
1.3877212802827528 previous
1.8225304898784933 emp.var.rate
2.348372356176547 blue-collar
1.14163524129405 entrepreneur
1.1512800634732347 housemaid
1.2940384604988566 management
1.4834937610887158 retired
1.1261108332839098 self-employed
1.4471899518339966 services
1.902603411017963 technician
1.3953800724789436 basic.6y
2.0079977473739965 basic.9y
2.6741950884656536 high.school
2.297931584308367 professional.course
3.3016071127980133 university.degree
1.6385893838857952 aug
1.0728676709603513 mar
2.2099886621689033 may
1.3450595960370224 nov
1.1412686812851054 oct
1.1068452368839075 sep
```

In []:

```
In [61]: logit = sm.Logit(Y.values, pc).fit()
display(logit.summary())
```

```
Optimization terminated successfully.
Current function value: 0.673024
Iterations 4
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

Logit Regression Results

```
=====
Dep. Variable:          y      No. Observations:      40949
Model:              Logit      Df Residuals:          40944
Method:              MLE       Df Model:              4
```

```

Date:          Fri, 10 May 2019   Pseudo R-squ.:      -0.9310
Time:          14:06:14          Log-Likelihood:     -27560.
converged:     True              LL-Null:             -14272.
                                   LLR p-value:         1.000

```

```

=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
x1              0.0352       0.010       3.449      0.001       0.015      0.055
x2             -0.1188       0.010     -11.713      0.000      -0.139     -0.099
x3              0.3641       0.010     35.551      0.000       0.344      0.384
x4             -0.1408       0.010     -13.967      0.000      -0.161     -0.121
x5             -0.0248       0.010      -2.464      0.014      -0.044     -0.005
=====

```

```
"""
```

```
In [ ]:
```

1. AdaBoost Classifier

```
In [62]: X_train, X_test, Y_train, Y_test = train_test_split(X.values, Y.values, test_size=0.9)
```

```

clf_rdfore = AdaBoostClassifier()
clf_rdfore = clf_rdfore.fit(X_train,Y_train)

```

```
y_hat = clf_rdfore.predict(X_test)
```

```

print('Test accuracy:',np.round(clf_rdfore.score(X_test, Y_test),3), ', Cross Entropy
Loss is:', log_loss(Y_test,y_hat))
calculate_metrics(Y_test,y_hat)
plot_ROC(Y_test,X_test,clf_rdfore)
MAE(Y_test,y_hat)

```


```
Test accuracy: 0.886 , Cross Entropy Loss is: 3.9276229223650616
```


```
Confusion matrix is:
```

```
[[32074  664]
 [ 3527  590]]
```

```
We have 32664 correct observations and 4191 misclassifications.
```

	precision	recall	f1-score	support
0	0.90	0.98	0.94	32738
1	0.47	0.14	0.22	4117
micro avg	0.89	0.89	0.89	36855
macro avg	0.69	0.56	0.58	36855
weighted avg	0.85	0.89	0.86	36855





Out[62]: 4191

2. SVC

VII UNBALANCED DATA PROBLEM

```
In [63]: sns.set_style('whitegrid')
```

```
In [64]: def make_balanced(X,Y):
    DF = X.copy()
    DF['Y'] = Y
    DF_Yes = DF[DF.Y==1]#.info()
    DF_No = DF[DF.Y==0]#.info()
    DF_B = DF_No.sample(len(DF_Yes))
    D = pd.concat([DF_B,DF_Yes])
    D = D.reindex(np.random.permutation(D.index))
```

```
D = D[:int(len(D)/2)]#, D[int(len(D)/2):]
return D
```

```
In [65]: D = make_balanced(X,Y)
```

```
In [66]: D.shape
```

```
Out[66]: (4544, 32)
```

```
In [67]: X_B = D.drop(columns='Y')
         Y_B = D['Y']
```

```
In [68]: X_train, X_test, Y_train, Y_test = train_test_split(X_B.values, Y_B.values,
        test_size=0.25)
        clf_rdfore = AdaBoostClassifier()
        clf_rdfore = clf_rdfore.fit(X_train,Y_train)

        y_hat = clf_rdfore.predict(X_test)

        print('Test accuracy:',np.round(clf_rdfore.score(X_test, Y_test),3), ', Cross Entropy
        Loss is:', log_loss(Y_test,y_hat))
        calculate_metrics(Y_test,y_hat)
        plot_ROC(Y_test,X_test,clf_rdfore)
        MAE(Y_test,y_hat)
```


Test accuracy: 0.729 , Cross Entropy Loss is: 9.364452017775967


Confusion matrix is:

```
[[477 93]
 [215 351]]
```

We have 828 correct observations and 308 misclassifications.

	precision	recall	f1-score	support
0	0.69	0.84	0.76	570
1	0.79	0.62	0.70	566
micro avg	0.73	0.73	0.73	1136
macro avg	0.74	0.73	0.73	1136
weighted avg	0.74	0.73	0.73	1136





Out[68]: 308

```
In [69]: X_R = X.drop(D.index)
         Y_R = Y.drop(D.index)
```

```
In [70]: C = [LogisticRegression(), AdaBoostClassifier(), KNeighborsClassifier(),
              RandomForestClassifier()]
```

```
def calc_sampling_imb(X,Y):
    L_NAME = ['lr','ada','knn','rforest']
    L_X, L_Y, L_YHAT, L_PROB = [],[],[],[]

    for cl in C:
        y_hat = cross_val_predict(cl, X, Y, cv=10, n_jobs=-1)
        probs = cross_val_predict(cl, X, Y, cv=10, method='predict_proba',n_jobs=-1)

        #print('Test accuracy:',np.round(c.score(X_R, Y_R),3), ', Cross Entropy Loss
is:', log_loss(Y_R,y_hat))
```

```

        calculate_metrics(Y,y_hat)
        #plot_ROC(Y_R,X,c)
        print('MAE:',MAE(Y,y_hat))

        #L_NAME.append(name)
        L_X.append(X.values)
        L_Y.append(Y.values)
        L_YHAT.append(y_hat)
        L_PROB.append(probs)

    plot_ROC2(L_Y,L_X,L_YHAT,L_PROB,L_NAME)

    #plot_ROC3(L_Y,L_X,L_CLASS,L_NAME)

```


```
In [71]: calc_sampling_imb(X_R,Y_R)
```

Confusion matrix is:

```
[[30064  4050]
 [ 2101   190]]
```


We have 30254 correct observations and 6151 misclassifications.

	precision	recall	f1-score	support
0	0.93	0.88	0.91	34114
1	0.04	0.08	0.06	2291
micro avg	0.83	0.83	0.83	36405
macro avg	0.49	0.48	0.48	36405
weighted avg	0.88	0.83	0.85	36405



```
MAE: 6151
Confusion matrix is:
[[32027  2087]
 [ 2131   160]]
We have 32187 correct observations and 4218 misclassifications.
```


	precision	recall	f1-score	support
0	0.94	0.94	0.94	34114
1	0.07	0.07	0.07	2291
micro avg	0.88	0.88	0.88	36405
macro avg	0.50	0.50	0.50	36405
weighted avg	0.88	0.88	0.88	36405



```
MAE: 4218
Confusion matrix is:
[[32049  2065]
 [ 2179   112]]
We have 32161 correct observations and 4244 misclassifications.
precision    recall  f1-score   support

0           0.94     0.94     0.94     34114
1           0.05     0.05     0.05      2291


micro avg    0.88     0.88     0.88    36405
macro avg    0.49     0.49     0.49    36405
weighted avg 0.88     0.88     0.88    36405
```


```
MAE: 4244
Confusion matrix is:
[[29448  4666]
 [ 2078   213]]
We have 29661 correct observations and 6744 misclassifications.
      precision    recall  f1-score   support

     0       0.93     0.86     0.90     34114
     1       0.04     0.09     0.06      2291

 micro avg       0.81     0.81     0.81    36405
 macro avg       0.49     0.48     0.48    36405
weighted avg       0.88     0.81     0.84    36405
```



MAE: 6744



```
In [72]: def train_B_test(X,Y):
    L_NAME = ['lr','ada','knn','rforest']
    L_X, L_Y, L_YHAT, L_PROB = [],[],[],[]

    D = make_balanced(X,Y)
    X_B = D.drop(columns='Y')
    Y_B = D['Y']

    X_R = X.drop(D.index)
    Y_R = Y.drop(D.index)

    for i,cl in enumerate(C):

        c = cl.fit(X_B,Y_B)
        y_hat = c.predict(X_R)
        probs = c.predict_proba(X_R)
        print(L_NAME[i])
        print('Test accuracy:',np.round(c.score(X_R, Y_R),3), ', Cross Entropy Loss
```

```

is:', log_loss(Y_R,y_hat))
    calculate_metrics(Y_R,y_hat)
    MAE(Y_R,y_hat)

    L_X.append(X_R.values)
    L_Y.append(Y_R.values)
    L_YHAT.append(y_hat)
    L_PROB.append(probs)

plot_ROC2(L_Y,L_X,L_YHAT,L_PROB,L_NAME)

```


```
In [73]: train_B_test(X,Y)
```

```

lr
Test accuracy: 0.797 , Cross Entropy Loss is: 6.995183247290505
Confusion matrix is:
[[27531  6563]
 [  810 1501]]
We have 29032 correct observations and 7373 misclassifications.

```

	precision	recall	f1-score	support
0	0.97	0.81	0.88	34094
1	0.19	0.65	0.29	2311
micro avg	0.80	0.80	0.80	36405
macro avg	0.58	0.73	0.59	36405
weighted avg	0.92	0.80	0.84	36405



ada

Test accuracy: 0.827 , Cross Entropy Loss is: 5.971470810230238


Confusion matrix is:

```
[[28671  5423]
```

```
 [  871 1440]]
```

We have 30111 correct observations and 6294 misclassifications.

	precision	recall	f1-score	support
0	0.97	0.84	0.90	34094
1	0.21	0.62	0.31	2311
micro avg	0.83	0.83	0.83	36405
macro avg	0.59	0.73	0.61	36405
weighted avg	0.92	0.83	0.86	36405



knn

Test accuracy: 0.71 , Cross Entropy Loss is: 9.999902631491398


Confusion matrix is:

```
[[24399  9695]
```

```
 [  845  1466]]
```

We have 25865 correct observations and 10540 misclassifications.


	precision	recall	f1-score	support
0	0.97	0.72	0.82	34094
1	0.13	0.63	0.22	2311
micro avg	0.71	0.71	0.71	36405
macro avg	0.55	0.67	0.52	36405
weighted avg	0.91	0.71	0.78	36405




```
rforest
Test accuracy: 0.737 , Cross Entropy Loss is: 9.069168676662667
Confusion matrix is:
[[25437  8657]
 [  902 1409]]
We have 26846 correct observations and 9559 misclassifications.
      precision    recall  f1-score   support

     0       0.97      0.75      0.84     34094
     1       0.14      0.61      0.23      2311

 micro avg       0.74      0.74      0.74     36405
 macro avg       0.55      0.68      0.53     36405
weighted avg       0.91      0.74      0.80     36405
```





VII Test Over/Under Sampling Methods

```
In [74]: sns.set_style('whitegrid')
```

```
In [96]: def calc_performance(c,name):

    classifier = [name, c]

    over_samplers = [
        #['ADA-SYN', ADASYN()],
        ['ROS', RandomOverSampler()],
        ['RUS', RandomUnderSampler()],
        #['SMOTENC', SMOTENC()],
        ['SMOTEENN', SMOTEENN()],
        ['ENN', EditedNearestNeighbours()],
    ]
```

```

pipelines = [
    ['{}-{}'.format(s[0], classifier[0]),
     make_pipeline(s[1], classifier[1])]
    for s in over_samplers
]

calc_sampling(pipelines,X,Y)

def calc_sampling(pipelines,X,Y):
    L_NAME, L_X, L_Y, L_YHAT, L_PROB = [],[],[],[],[]

    X_train, X_test, Y_train, Y_test = train_test_split(X.values, Y.values,
test_size=0.85)
    for name, pipeline in pipelines:

        p = pipeline.fit(X_train,Y_train)
        y_hat = p.predict(X_test)
        probs = p.predict_proba(X_test)
        calculate_metrics(Y_test,y_hat)

        L_NAME.append(name)
        L_X.append(X_test)
        L_Y.append(Y_test)
        L_YHAT.append(y_hat)
        L_PROB.append(probs)

    plot_ROC2(L_Y,L_X,L_YHAT,L_PROB,L_NAME)

```

```

In [97]: NAME = ['LogisticRegression','AdaBoost','KNN','RForest']
         C_dict = dict(zip(NAME,C))

         for name,c in C_dict.items():
             calc_performance(c,name)

```

Confusion matrix is:


```

[[24535  6421]
 [ 1372  2479]]

```

We have 27014 correct observations and 7793 misclassifications.

	precision	recall	f1-score	support
0	0.95	0.79	0.86	30956
1	0.28	0.64	0.39	3851
micro avg	0.78	0.78	0.78	34807
macro avg	0.61	0.72	0.63	34807
weighted avg	0.87	0.78	0.81	34807




Confusion matrix is:

```
[[23994  6962]
```

```
 [ 1293  2558]]
```


We have 26552 correct observations and 8255 misclassifications.

	precision	recall	f1-score	support
0	0.95	0.78	0.85	30956
1	0.27	0.66	0.38	3851
micro avg	0.76	0.76	0.76	34807
macro avg	0.61	0.72	0.62	34807
weighted avg	0.87	0.76	0.80	34807




Confusion matrix is:
[[19637 11319]
[969 2882]]
We have 22519 correct observations and 12288 misclassifications.


	precision	recall	f1-score	support
0	0.95	0.63	0.76	30956
1	0.20	0.75	0.32	3851
micro avg	0.65	0.65	0.65	34807
macro avg	0.58	0.69	0.54	34807
weighted avg	0.87	0.65	0.71	34807



Confusion matrix is:
[[28834 2122]
 [2218 1633]]
We have 30467 correct observations and 4340 misclassifications.

	precision	recall	f1-score	support
0	0.93	0.93	0.93	30956
1	0.43	0.42	0.43	3851
micro avg	0.88	0.88	0.88	34807
macro avg	0.68	0.68	0.68	34807
weighted avg	0.87	0.88	0.87	34807






Confusion matrix is:

```
[[24417 6545]
 [ 1323 2522]]
```


We have 26939 correct observations and 7868 misclassifications.

	precision	recall	f1-score	support
0	0.95	0.79	0.86	30962
1	0.28	0.66	0.39	3845
micro avg	0.77	0.77	0.77	34807
macro avg	0.61	0.72	0.63	34807
weighted avg	0.87	0.77	0.81	34807




Confusion matrix is:
[[23937 7025]
 [1380 2465]]
We have 26402 correct observations and 8405 misclassifications.

	precision	recall	f1-score	support
0	0.95	0.77	0.85	30962
1	0.26	0.64	0.37	3845
micro avg	0.76	0.76	0.76	34807
macro avg	0.60	0.71	0.61	34807
weighted avg	0.87	0.76	0.80	34807



```
Confusion matrix is:
[[25787  5175]
 [ 1521  2324]]
We have 28111 correct observations and 6696 misclassifications.
```


	precision	recall	f1-score	support
0	0.94	0.83	0.89	30962
1	0.31	0.60	0.41	3845
micro avg	0.81	0.81	0.81	34807
macro avg	0.63	0.72	0.65	34807
weighted avg	0.87	0.81	0.83	34807




```
Confusion matrix is:
[[28720  2242]
 [ 2162 1683]]
We have 30403 correct observations and 4404 misclassifications.
      precision    recall  f1-score   support

     0       0.93      0.93      0.93     30962
     1       0.43      0.44      0.43      3845


 micro avg       0.87      0.87      0.87     34807
 macro avg       0.68      0.68      0.68     34807
weighted avg       0.87      0.87      0.87     34807
```





Confusion matrix is:
[[23602 7356]
 [1930 1919]]
We have 25521 correct observations and 9286 misclassifications.

	precision	recall	f1-score	support
0	0.92	0.76	0.84	30958
1	0.21	0.50	0.29	3849
micro avg	0.73	0.73	0.73	34807
macro avg	0.57	0.63	0.56	34807
weighted avg	0.85	0.73	0.78	34807




Confusion matrix is:

```
[[21674  9284]
 [ 1446  2403]]
```


We have 24077 correct observations and 10730 misclassifications.

	precision	recall	f1-score	support
0	0.94	0.70	0.80	30958
1	0.21	0.62	0.31	3849
micro avg	0.69	0.69	0.69	34807
macro avg	0.57	0.66	0.56	34807
weighted avg	0.86	0.69	0.75	34807




Confusion matrix is:
[[18137 12821]
 [1107 2742]]
We have 20879 correct observations and 13928 misclassifications.


	precision	recall	f1-score	support
0	0.94	0.59	0.72	30958
1	0.18	0.71	0.28	3849
micro avg	0.60	0.60	0.60	34807
macro avg	0.56	0.65	0.50	34807
weighted avg	0.86	0.60	0.67	34807



Confusion matrix is:
[[28798 2160]
 [2626 1223]]
We have 30021 correct observations and 4786 misclassifications.

	precision	recall	f1-score	support
0	0.92	0.93	0.92	30958
1	0.36	0.32	0.34	3849
micro avg	0.86	0.86	0.86	34807
macro avg	0.64	0.62	0.63	34807
weighted avg	0.86	0.86	0.86	34807






Confusion matrix is:

```
[[29043 1899]
 [ 2776 1089]]
```

We have 30132 correct observations and 4675 misclassifications.


	precision	recall	f1-score	support
0	0.91	0.94	0.93	30942
1	0.36	0.28	0.32	3865
micro avg	0.87	0.87	0.87	34807
macro avg	0.64	0.61	0.62	34807
weighted avg	0.85	0.87	0.86	34807



```
Confusion matrix is:
[[22726  8216]
 [ 1438  2427]]
We have 25153 correct observations and 9654 misclassifications.
      precision    recall  f1-score   support


     0       0.94      0.73      0.82     30942
     1       0.23      0.63      0.33      3865

 micro avg       0.72      0.72      0.72     34807
 macro avg       0.58      0.68      0.58     34807
weighted avg       0.86      0.72      0.77     34807
```




Confusion matrix is:
[[27213 3729]
 [1939 1926]]
We have 29139 correct observations and 5668 misclassifications.


	precision	recall	f1-score	support
0	0.93	0.88	0.91	30942
1	0.34	0.50	0.40	3865
micro avg	0.84	0.84	0.84	34807
macro avg	0.64	0.69	0.66	34807
weighted avg	0.87	0.84	0.85	34807



```
Confusion matrix is:
[[28494  2448]
 [ 2278 1587]]
We have 30081 correct observations and 4726 misclassifications.
```

	precision	recall	f1-score	support
0	0.93	0.92	0.92	30942
1	0.39	0.41	0.40	3865
micro avg	0.86	0.86	0.86	34807
macro avg	0.66	0.67	0.66	34807
weighted avg	0.87	0.86	0.87	34807





VII selection of relevant variables from balanced dataset

```
In [86]: L, Counts, Coeffs = [], [], []
         for i in range(1000):
             D = make_balanced(X, Y)
             X_B = D.drop(columns='Y')
             Y_B = D['Y']

             X_B = sm.add_constant(X_B)

             logit = sm.Logit(Y_B, X_B).fit(dis=0)
             #display(logit.summary())

             alpha = 0.10
             a = logit.pvalues < alpha
```

```


X4 = X_B[X_B.columns[a]]
#print("Not Statistically significant regressors are:")
temp = list(X_B.columns[a])
L.append(temp)
Counts.append(len(temp))
Coeffs.append(logit.params.loc[temp])

```

In []:

In [87]: L = [y for x in L for y in x]

In [88]: plt.figure(figsize=(12,8))
plt.hist(L,bins=len(set(L)))
plt.xticks(rotation=90)
plt.show()



In [89]: from collections import Counter
d = Counter(L)

```

import operator
sorted_d = dict(sorted(d.items(), key=operator.itemgetter(1), reverse=True))

```

In [90]: TopVars = []
print('Most frequently significant variables:')
for i, (k,v) in enumerate(sorted_d.items()):
 if v > 500:

```
print(k, ': ', v)
TopVars.append(k)
```

Most frequently significant variables:

```
previous : 1000
emp.var.rate : 1000
mar : 1000
may : 1000
oct : 974
sep : 972
campaign : 954
jul : 940
nov : 822
student : 702
retired : 683
```

```
In [91]: V = []
        for v in TopVars:
            summ = 0
            count = 0
            for c in Coeffs:
                try:
                    summ += c.loc[v]
                    count += 1
                except KeyError:
                    pass
            V.append(summ/count)
```

```
In [92]: list(np.round(np.array(V),3))
```

```
Out[92]: [0.348, -0.481, 1.32, -0.601, 0.881, 0.88, -0.049, 0.495, -0.425, 0.578, 0.493]
```

```
In [93]: list(np.round(np.exp(np.array(V)),3))
```

```
Out[93]: [1.417, 0.618, 3.742, 0.548, 2.413, 2.412, 0.952, 1.641, 0.654, 1.782, 1.637]
```

```
In [ ]:
```